

Scalable Machine Learning and Cloud Computing

Xianyuan Liu

xianyuan.liu@sheffield.ac.uk

HPC = High-performance computing cluster
AWS = Amazon Web Services

In Today's Lecture:

Today we will,

- introduce local, HPC and Cloud Computing.
- look at **whether** to distribute a given problem
- have a quick run through of **services** and **tools** from the field of Cloud Computing.
- have several **walk-throughs** (AWS **EC2** and **lambda**).

In the Lab on Friday:

On Friday you will,

- **create** and **connect** to an **instance** on AWS
- set up an **Elastic Map Reduce cluster**
- use it to run **Apache Spark** to solve a matrix factorisation problem
- set up a **Lambda function**

HPC = High-performance computing cluster

Local vs HPC vs Cloud Computing

- Local Computing
 - Physically and directly managed by users, normally in the same geographic location as the user.

HPC = High-performance computing cluster

Local vs HPC vs Cloud Computing

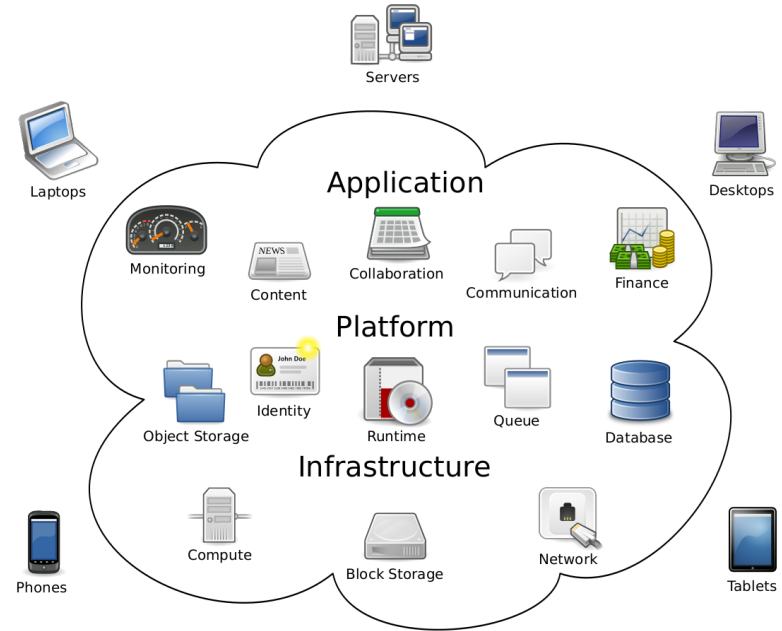
- Local Computing
 - Physically and directly managed by users, normally in the same geographic location as the user.
- HPC
 - Clusters of computers combined to form a powerful computing environment (Stanage - Manchester, Bessemer - Leeds)
- Cloud Computing



The old Iceberg system server racks.

What is Cloud Computing

- Cloud computing is the **on-demand** availability of computer system resources, without direct active management by the user.
- Cloud computing relies on the **sharing of resources** to achieve coherence and typically uses a **pay-as-you-go** model.
- Large clouds often have functions **distributed** over multiple locations.



HPC vs Cloud?

Until now we've used the university's own HPC system.

HPC:

- Often provided for **free** by institution
- highly **interconnected** nodes - vital for many problems
- Most HPCs use SLURM (includes tools like job array, which let you run embarrassingly parallel arrays of jobs).
- Provide **support**



HPC vs Cloud?

HPC:

- Often provided for **free** by institution
- highly **interconnected** nodes - vital for many problems
- Most HPCs use SLURM (includes tools like job array, which let you run embarrassingly parallel arrays of jobs).
- Provide **support**

Cloud:

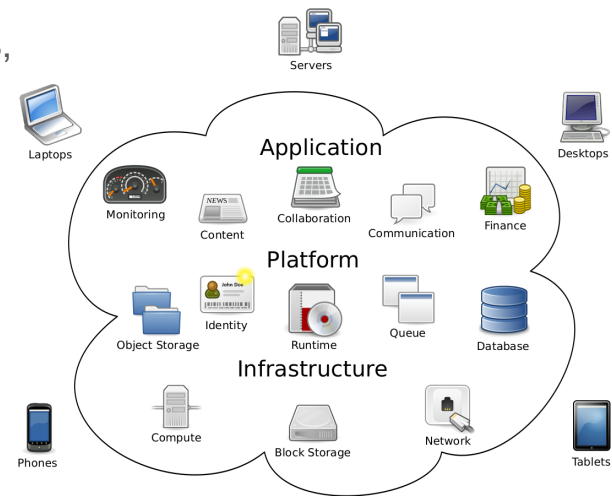
- **No queuing**
- **Quick** to set up
- Often **only option** outside of universities
- Well suited for **embarrassingly parallel** problems
- Almost all **platforms** and features **supported**
- Appropriate if **hosting a front-end server**
- Can be very expensive (+ accidental spend!)



Often 'rent' (virtual) hardware while you need it. Large clusters can be created briefly.

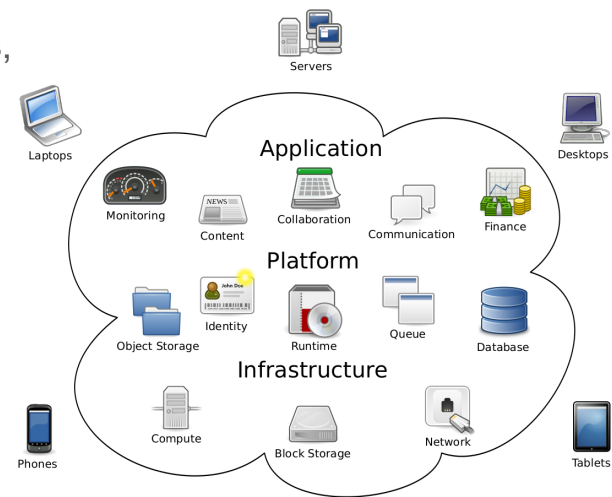
Why Cloud Computing for Machine Learning

- Powerful Hardware
 - **On-demand** access to state-of-the-art hardware resources, instead of expensive investments in physical hardware.



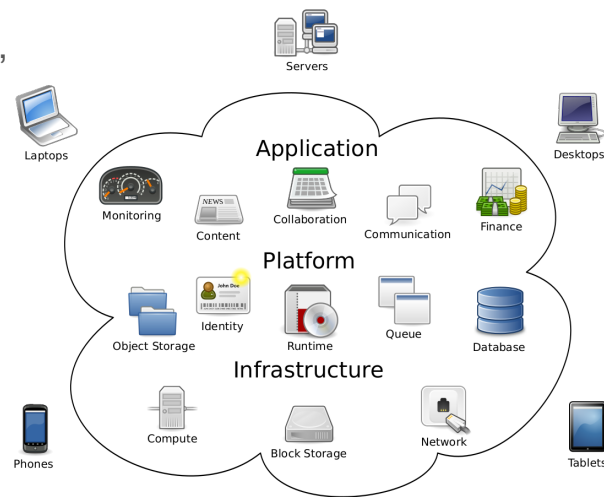
Why Cloud Computing for Machine Learning

- Powerful Hardware
 - **On-demand** access to state-of-the-art hardware resources, instead of expensive investments in physical hardware.
- Scalability
 - **Dynamically** adjust resources (e.g., hardware, memory, storage) as your needs change.



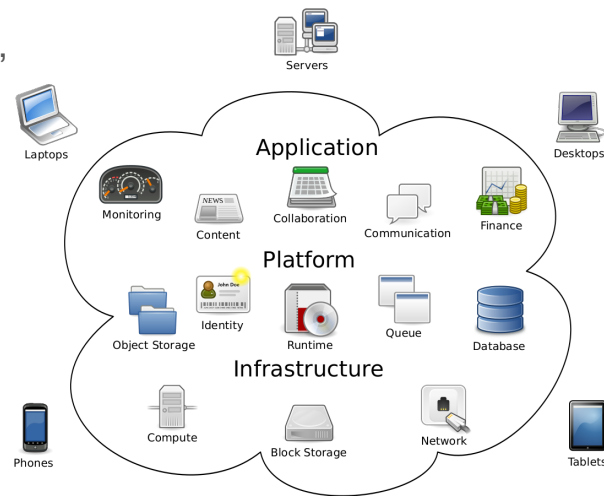
Why Cloud Computing for Machine Learning

- Powerful Hardware
 - **On-demand** access to state-of-the-art hardware resources, instead of expensive investments in physical hardware.
- Scalability
 - **Dynamically** adjust resources (e.g., hardware, memory, storage) as your needs change.
- Pay-as-you-go Model
 - Pay **only** for the compute time and storage used.
 - Achieve significant savings on energy consumption, facility management, and ongoing hardware maintenance.



Why Cloud Computing for Machine Learning

- Powerful Hardware
 - **On-demand** access to state-of-the-art hardware resources, instead of expensive investments in physical hardware.
- Scalability
 - **Dynamically** adjust resources (e.g., hardware, memory, storage) as your needs change.
- Pay-as-you-go Model
 - Pay **only** for the compute time and storage used.
 - Achieve significant savings on energy consumption, facility management, and ongoing hardware maintenance.
- Flexibility
 - Access from **anywhere** with an internet connection.
 - Share resources with **multiple** users.



Local vs HPC vs Cloud Computing

- Local Computing
 - Physically and directly managed by users, normally in the same geographic location as the user.

- HPC
 - Clusters of computers combined to form a powerful computing environment (Stanage - Manchester, Bessemer - Leeds)

- Cloud Computing

Distributed system - Parallel computing





The old Iceberg system server racks.

Distributed system - Parallel computing

Different levels of parallelism:

- **Single process** (e.g. python GIL code)
- Well parallelised multiprocessing code (might be good to deploy on a single instance with lots of cores on AWS or HPC)
- Distributed over **multiple nodes** (either on AWS or HPC).




...



Distributed system - Parallel computing

Different levels of parallelism:

- **Single process** (e.g. python GIL code)
- Well parallelised multiprocessing code (might be good to deploy on a single instance with lots of cores on AWS or HPC)
- Distributed over **multiple nodes** (either on AWS or HPC).



Local vs Distributed

- It's not always worth parallelising your application!

Local vs Distributed

- It's not always worth parallelising your application!
- **Human time** is worth more than **computer time**



Local vs Distributed

- It's not always worth parallelising your application!
- **Human time** is worth more than **computer time**, so if you can do something else and be patient it might be best not to bother with rewriting it to be parallel.

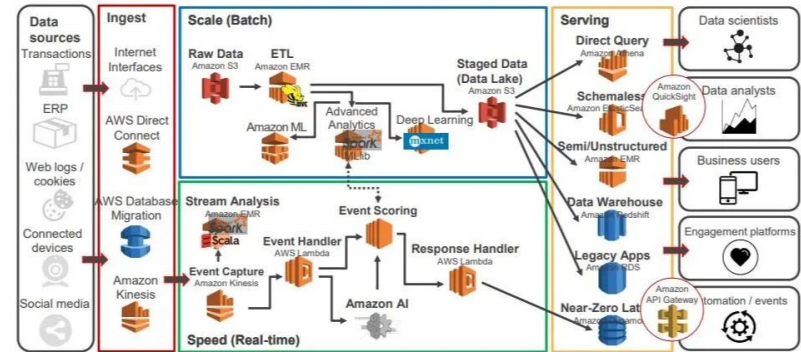


Why stay Local...

- May need to **redesign** algorithm

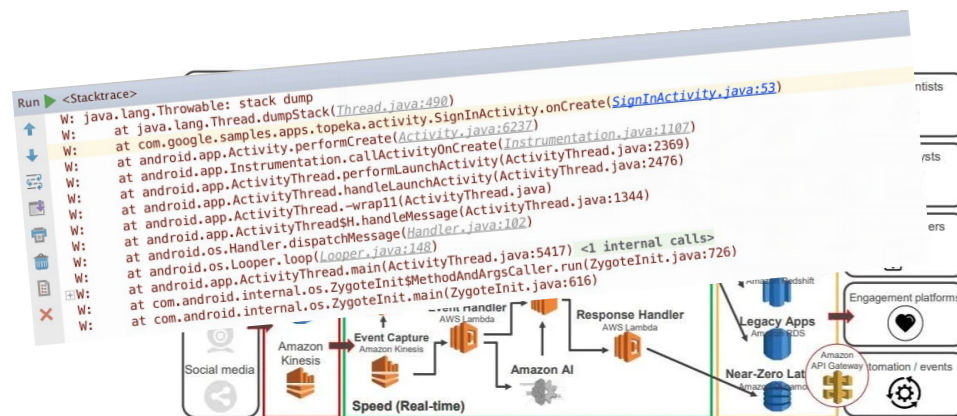
Why stay Local...

- May need to **redesign** algorithm
- Managing the **pipeline**.



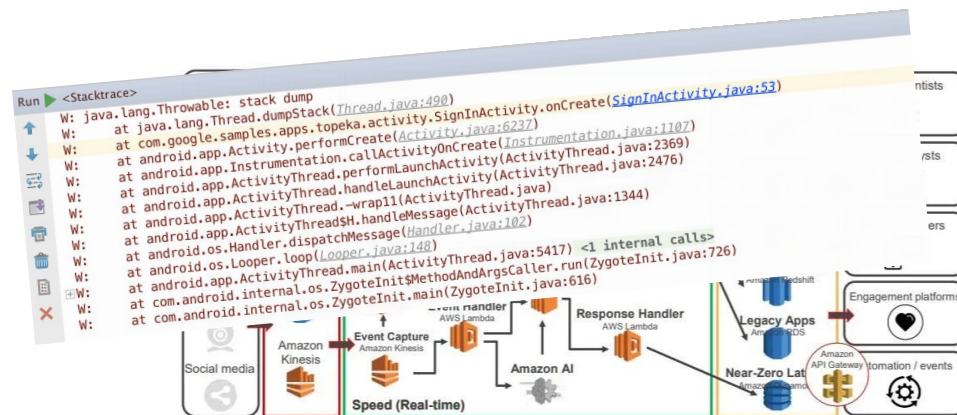
Why stay Local...

- May need to **redesign** algorithm
- Managing the **pipeline**.
- Typically such environments are less **reliable** and less **predictable**.
 - Far more difficult to **debug**.



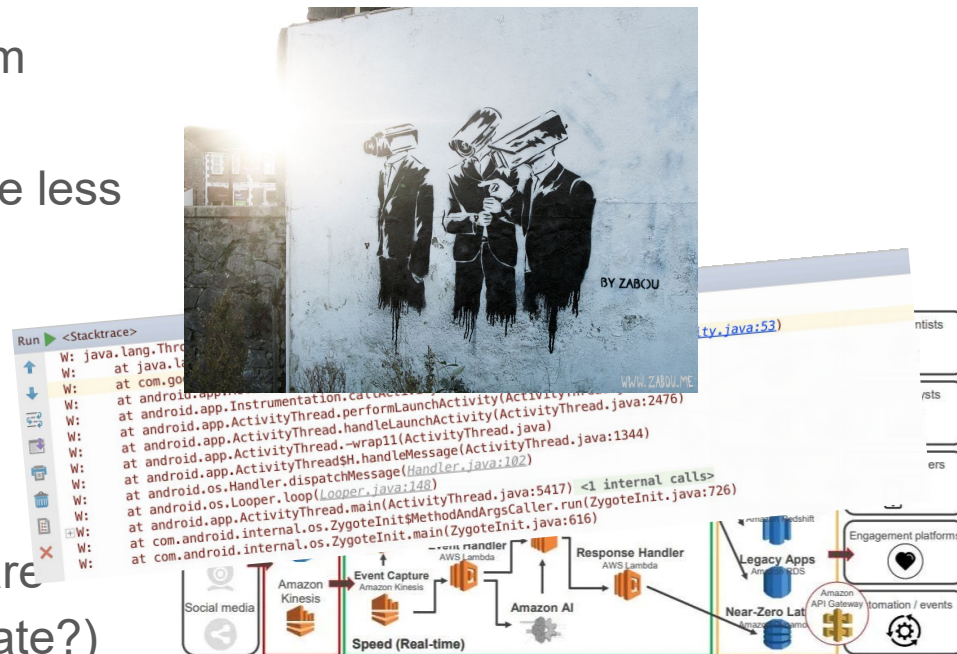
Why stay Local...

- May need to **redesign** algorithm
- Managing the **pipeline**.
- Typically such environments are less **reliable** and less **predictable**.
 - Far more difficult to **debug**.
- **Shifting data** to/from the servers might be a bottleneck.



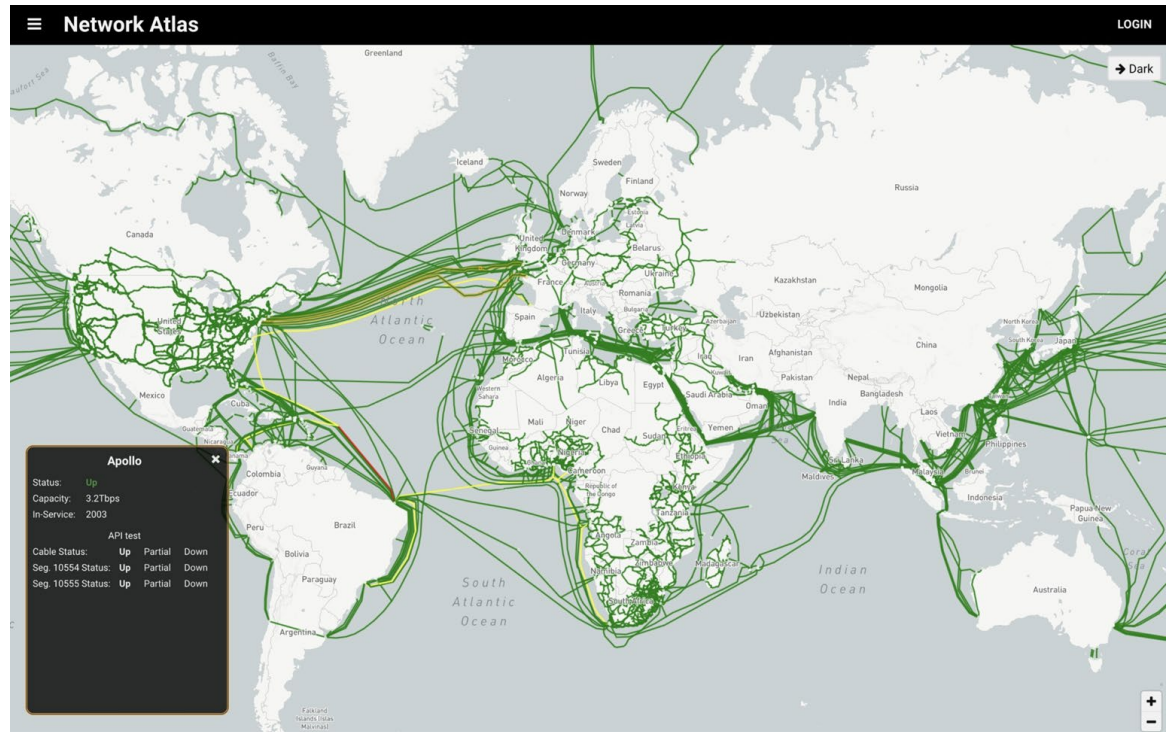
Why stay Local...

- May need to **redesign** algorithm
- Managing the **pipeline**.
- Typically such environments are less **reliable** and less **predictable**.
 - Far more difficult to **debug**.
- **Shifting data** to/from the servers might be a bottleneck.
- **Security** and **privacy** (where are you sending the data? Is it private?)



Privacy and the GDPR...

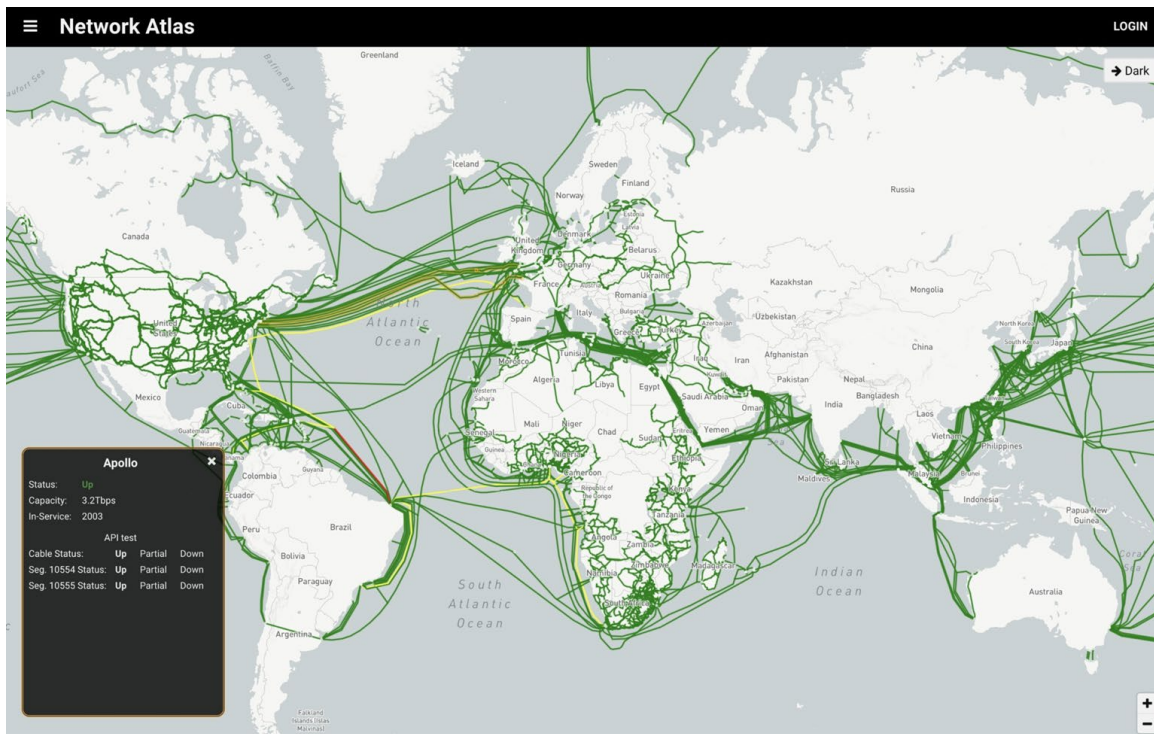
The **EU GDPR** (General Data Protection Regulation) restricts the transfer of **personal data** to countries outside of the EU.



Privacy and the GDPR...

Post-brexite: UK GDPR is now in law,
and is basically the same as the EU
GDPR for now.

The **EU GDPR** (General Data Protection Regulation)
restricts the transfer of
personal data to
countries outside of
the EU.



Privacy and the GDPR...

Post-brexite: UK GDPR is now in law,
and is basically the same as the EU
GDPR for now.

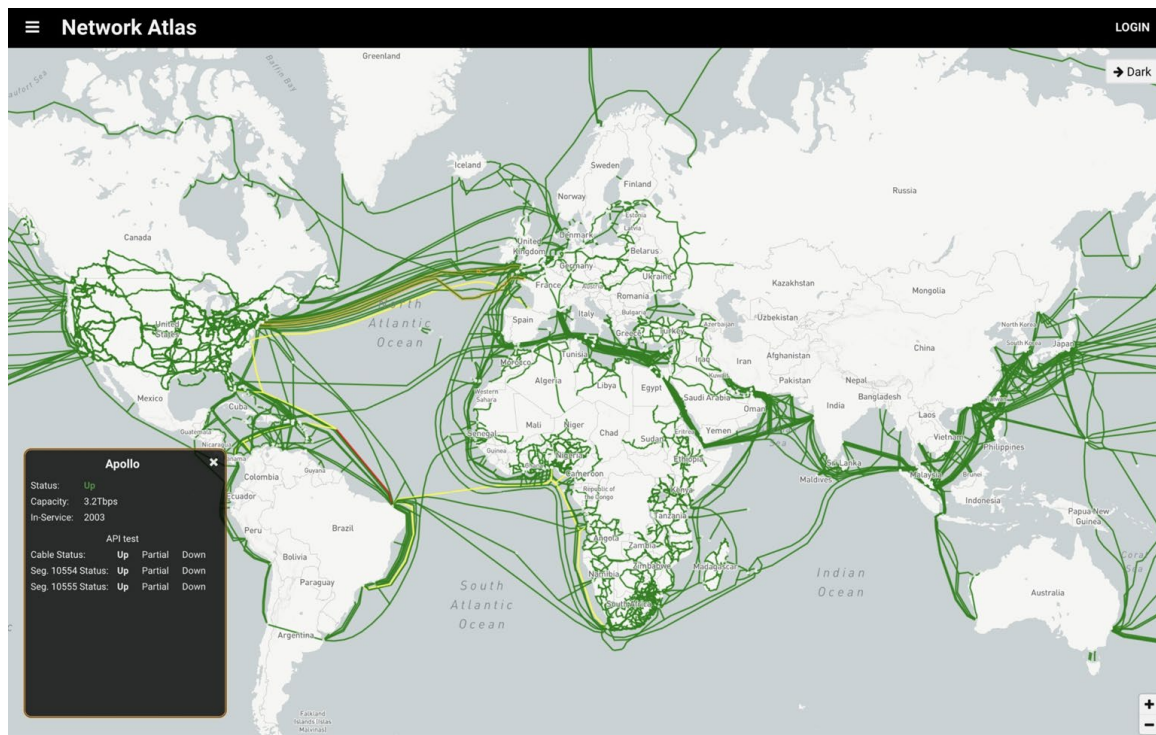
The **EU GDPR** (General Data Protection Regulation)

restricts the transfer of
personal data to
countries outside of
the EU.

One can process the data if it's
done in a country covered by an
EU/UK “**adequacy decision**”

The UK and EU consider each
other to have this.

Otherwise one needs a **standard
contractual clause** (these have
begun to diverge between the EU
and UK).



Privacy and the GDPR...

Post-brexite: UK GDPR is now in law,
and is basically the same as the EU
GDPR for now.

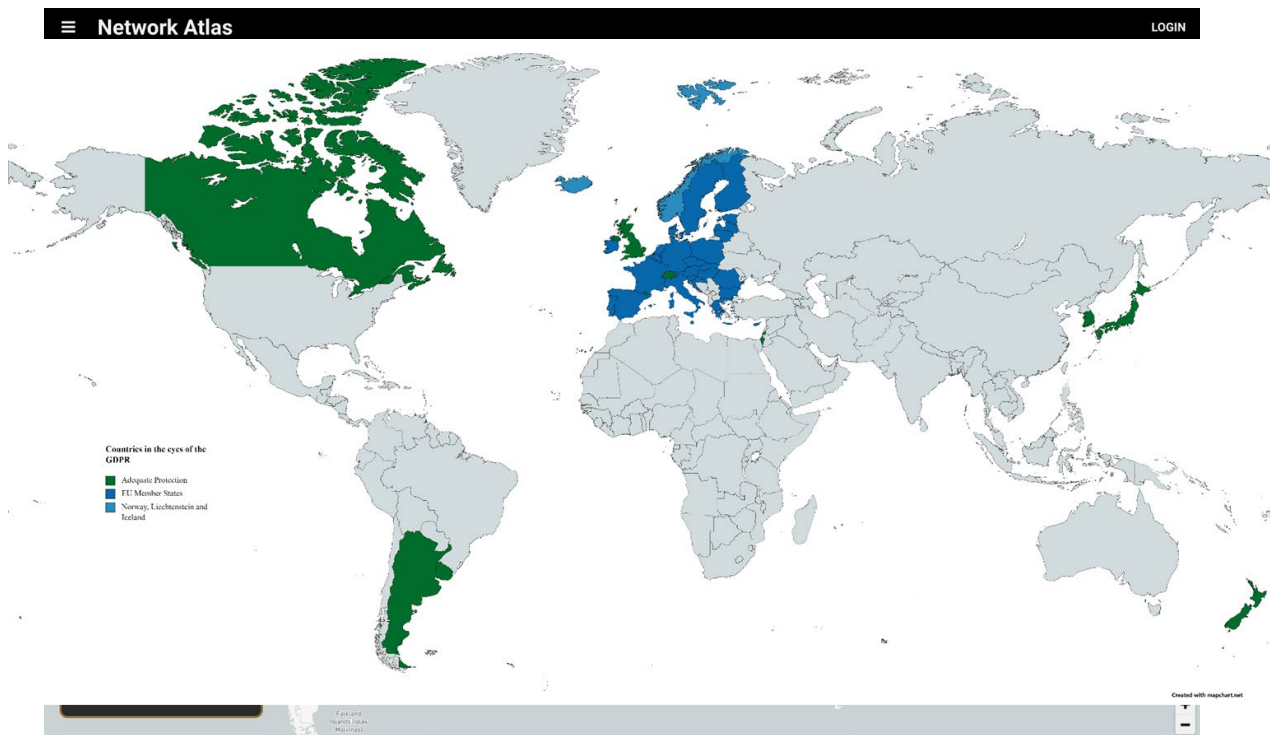
The **EU GDPR** (General Data Protection Regulation)

restricts the transfer of
personal data to
countries outside of
the EU.

One can process the data if it's
done in a country covered by an
EU/UK “**adequacy decision**”

The UK and EU consider each
other to have this.

Otherwise one needs a **standard
contractual clause** (these have
begun to diverge between the EU
and UK).



Privacy and the GDPR...

Post-brexite: UK GDPR is now in law,
and is basically the same as the EU
GDPR for now.

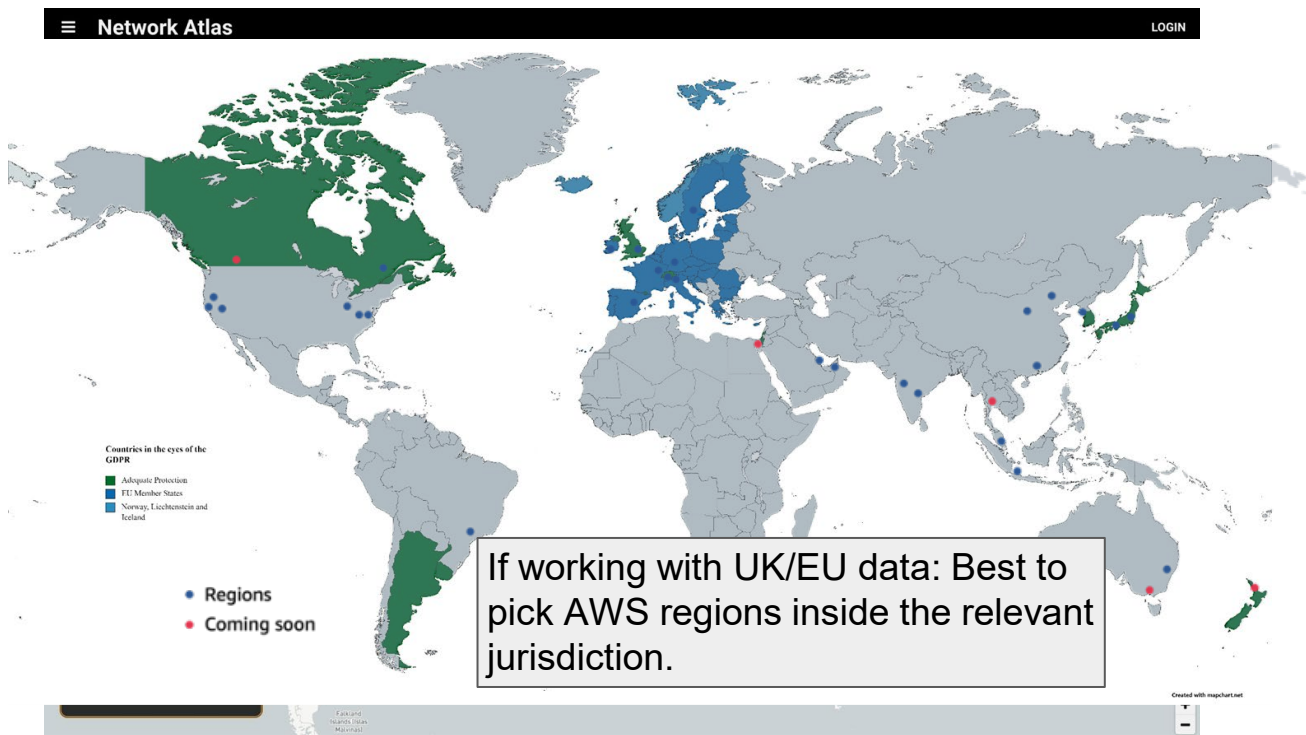
The **EU GDPR** (General Data Protection Regulation)

restricts the transfer of
personal data to
countries outside of
the EU.

One can process the data if it's
done in a country covered by an
EU/UK “**adequacy decision**”

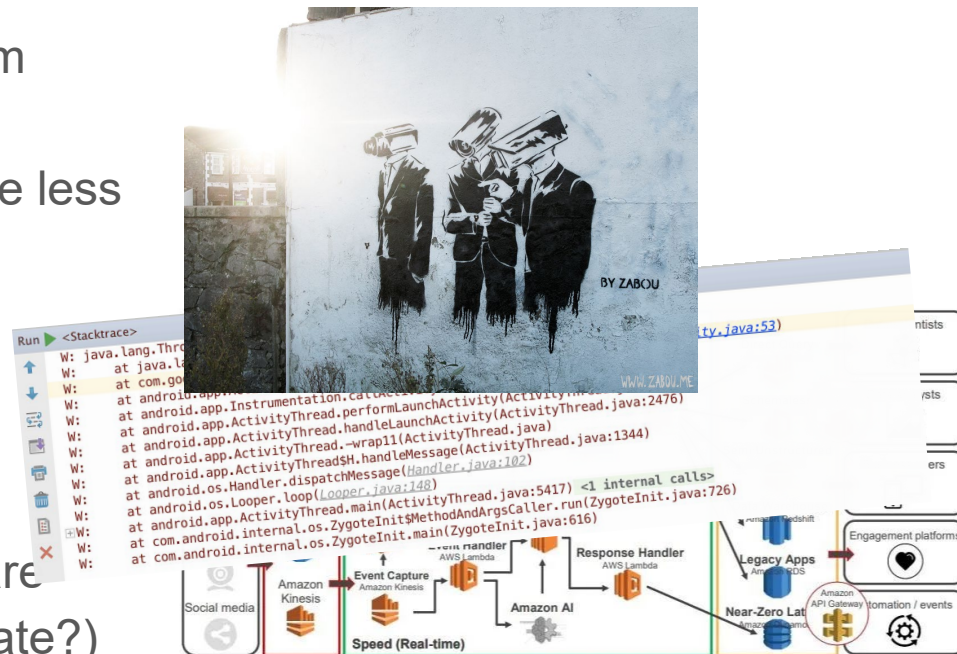
The UK and EU consider each
other to have this.

Otherwise one needs a **standard
contractual clause** (these have
begun to diverge between the EU
and UK).



Why stay Local...

- May need to **redesign** algorithm
- Managing the **pipeline**.
- Typically such environments are less **reliable** and less **predictable**.
 - Far more difficult to **debug**.
- **Shifting data** to/from the servers might be a bottleneck.
- **Security** and **privacy** (where are you sending the data? Is it private?)
- **Cheaper**.



Service Models

There are different **service models** available: the trade off is typically complexity vs flexibility:

- **Infrastructure** as a Service (IaaS) - high level of control
 - Provides virtual servers, storage, networking resources, and other essential elements.
 - E.g. **AWS EC2**, Microsoft Azure Virtual Machines, Google Compute Engine.

Service Models

There are different **service models** available: the trade off is typically complexity vs flexibility:

- **Infrastructure** as a Service (IaaS) - high level of control
 - Provides virtual servers, storage, networking resources, and other essential elements.
 - E.g. **AWS EC2**, Microsoft Azure Virtual Machines, Google Compute Engine.
- **Platform** as a Service (PaaS) [always on] - medium
- **Function** as a Service (FaaS) [event driven] - low
 - Offers a development and deployment platform on top of the infrastructure.
 - You provide software/code, but the service starts up virtual machines, allocates storage, runs databases, etc.
 - E.g. website with database / **AWS Lambda**, Azure Functions, Google Cloud Functions.

Service Models

There are different **service models** available: the trade off is typically complexity vs flexibility:

- **Infrastructure** as a Service (IaaS) - high level of control
 - Provides virtual servers, storage, networking resources, and other essential elements.
 - E.g. **AWS EC2**, Microsoft Azure Virtual Machines, Google Compute Engine.
- **Platform** as a Service (PaaS) [always on] - medium
- **Function** as a Service (FaaS) [event driven] - low
 - Offers a development and deployment platform on top of the infrastructure.
 - You provide software/code, but the service starts up virtual machines, allocates storage, runs databases, etc.
 - E.g. website with database / **AWS Lambda**, Azure Functions, Google Cloud Functions.
- **Software** as a Service (SaaS) - low
 - Provides ready-to-use software applications over the internet
 - E.g. Google Docs

Service Models

There are different **service models** available: the trade off is typically complexity vs flexibility:

- **Infrastructure** as a Service (IaaS) - high level of control
 - Provides virtual servers, storage, networking resources, and other essential elements.
 - E.g. **AWS EC2**, Microsoft Azure Virtual Machines, Google Compute Engine.
- **Platform** as a Service (PaaS) [always on] - medium
- **Function** as a Service (FaaS) [event driven] - low
 - Offers a development and deployment platform on top of the infrastructure.
 - You provide software/code, but the service starts up virtual machines, allocates storage, runs databases, etc.
 - E.g. website with database / **AWS Lambda**, Azure Functions, Google Cloud Functions.
- **Software** as a Service (SaaS) - low
 - Provides ready-to-use software applications over the internet
 - E.g. Google Docs


Infrastructure as Code (IaC) - cloud
configuration specified in code.

Cloud Computing Platforms

We will focus on **AWS** (others similar).

Amazon Maintains Cloud Lead as Microsoft Edges Closer

Worldwide market share of leading cloud infrastructure service providers in Q4 2023*



Cloud infrastructure service revenues in Q4 2023

\$73.7B

* Includes platform as a service (PaaS) and infrastructure as a service (IaaS) as well as hosted private cloud services

Source: Synergy Research Group



statista

Cloud Computing Platforms

We will focus on **AWS** (others similar).

Huge range of services.

Amazon Elastic
Compute Cloud (**EC2**)
allows users to rent
virtual machines.

Amazon Maintains Cloud Lead as Microsoft Edges Closer

Worldwide market share of leading cloud infrastructure
service providers in Q4 2022*

All services

Services by category



Compute

- EC2
- Lightsail
- Lambda
- Batch
- Elastic Beanstalk
- Serverless Application Repository
- AWS Outposts
- EC2 Image Builder
- AWS App Runner
- AWS SimSpace Weaver



Containers

- Elastic Container Registry
- Elastic Container Service
- Elastic Kubernetes Service
- Red Hat OpenShift Service on AWS



Storage



Management & Governance

- AWS Organizations
- CloudWatch
- AWS Auto Scaling
- CloudFormation
- AWS Config
- OpsWorks
- Service Catalog
- Systems Manager
- Trusted Advisor
- Control Tower
- AWS Well-Architected Tool
- AWS Chatbot
- Launch Wizard
- AWS Compute Optimizer
- Resource Groups & Tag Editor
- Amazon Grafana
- Amazon Prometheus
- AWS Resilience Hub
- Incident Manager



Security, Identity, & Compliance

- Resource Access Manager
- Cognito
- Secrets Manager
- GuardDuty
- Amazon Inspector
- Amazon Macie
- IAM Identity Center
- Certificate Manager
- Key Management Service
- CloudHSM
- Directory Service
- WAF & Shield
- AWS Firewall Manager
- AWS Artifact
- Detective
- AWS Signer
- AWS Private Certificate Authority
- Security Hub
- AWS Audit Manager

AWS: Examples of Instance Types

On-demand (pay per hour) or Reserved

AWS: Examples of Instance Types

On-demand (pay per hour) or Reserved

Cheapest: t4g.nano [burstable, smallest instance] for 3 years \$41 (ex. VAT).

- Also need to pay a few pence for storage, data transfer, etc.

AWS: Examples of Instance Types

On-demand (pay per hour) or Reserved

Cheapest: t4g.nano [burstable, smallest instance] for 3 years \$41 (ex. VAT).

- Also need to pay a few pence for storage, data transfer, etc.

Compute optimised

Instance name ▼	On-Demand hourly rate ▼	vCPU ▼	Memory ▼	Storage ▼	Network performance ▼
c7a.metal-48xl	\$9.85344	192	384 GiB	EBS Only	50000 Megabit

Memory optimised

Instance name ▼	On-Demand hourly rate ▼	vCPU ▼	Memory ▼	Storage ▼	Network performance ▼
u-24tb1.112xlarge	\$261.73	448	24576 GiB	EBS Only	100 Gigabit

AWS: Examples of Instance Types

Other decisions: Storage type (e.g. HDD vs SSD), what OS will you install, etc. We will explore this in more detail in the lab.

On-demand (pay per hour) or Reserved

Cheapest: t4g.nano [burstable, smallest instance] for 3 years \$41 (ex. VAT).

- Also need to pay a few pence for storage, data transfer, etc.

Compute optimised

Instance name ▼	On-Demand hourly rate ▼	vCPU ▼	Memory ▼	Storage ▼	Network performance ▼
c7a.metal-48xl	\$9.85344	192	384 GiB	EBS Only	50000 Megabit

Memory optimised

Instance name ▼	On-Demand hourly rate ▼	vCPU ▼	Memory ▼	Storage ▼	Network performance ▼
u-24tb1.112xlarge	\$261.73	448	24576 GiB	EBS Only	100 Gigabit

There are also instances that have one or more **GPUs**.

Some tools widely used...

- MapReduce



- Apache Spark



- Terraform



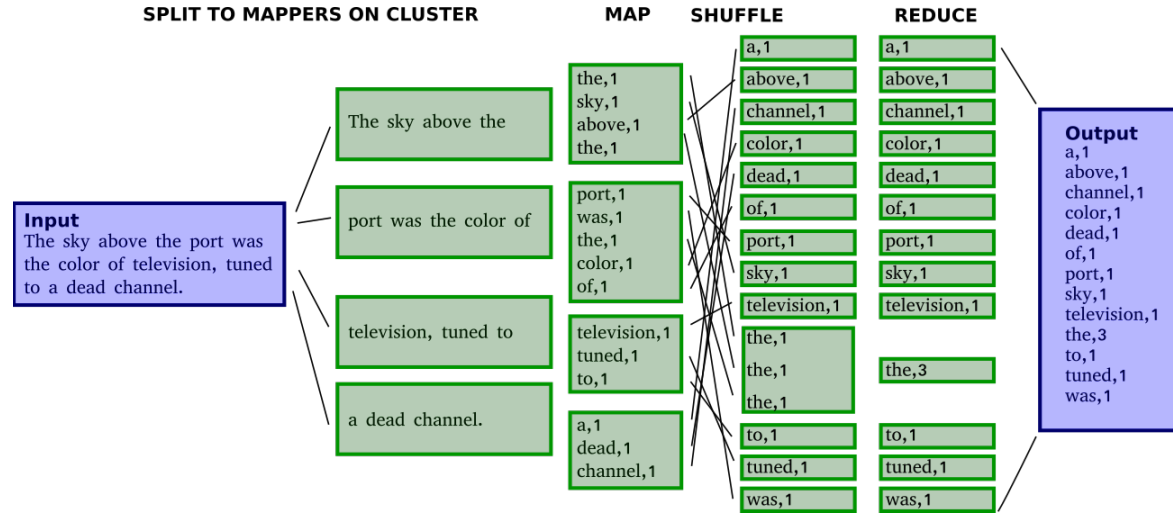
- Docker



docker

MapReduce

- A programming model for processing **large datasets** in parallel across clusters.
- First, **breaks down** a large task into smaller, independent subtasks (**map** phase)
- **Shuffles** the results to be in the right locations and then aggregates them (**reduce** phase).



Example of a distributed word-count, using MapReduce

Apache Spark



- Apache SPARK is a **distributed cluster** computing framework.
- More flexible and faster than MapReduce
 - Functionalities for batch data processing (similar to MapReduce), real-time data processing, and machine learning.
 - In-memory processing rather than disk-based processing.
- Particularly useful for **iterative algorithms** (e.g. ML gradient descent).

We'll look at this in the lab in more detail.

Spark's RDDs function as a [working set](#) for distributed programs that offers a (deliberately) restricted form of distributed [shared memory](#).

Inside Apache Spark the workflow is managed as a [directed acyclic graph](#) (DAG). Nodes represent RDDs while edges represent the operations on the RDDs.

- from wikipedia's [article](#).

Terraform



- Defines the cloud infrastructure (virtual servers, storage, networking) as code files and can automate the creation and management across different cloud providers.
- Essential to manage in a scalable and automated way, especially for complex deployments.
- **Infrastructure** as code (IaC)

<https://docs.aws.amazon.com/cli/latest/userguide/getting-started-install.html>
<https://app.terraform.io/app/getting-started/example>

Docker



- A platform allows to put all the **software, libraries and configuration** files that you need to run an application into a **container**.
- It sort of replaces **virtual machines**, in a 'lighter' way, and makes it easier to **move data** between them.
- Simplifies development workflows, and ensures consistent application behavior across different environments.
- Mainly for linux, but typically not on HPC.

Other service... Lambda



AWS Lambda

- AWS Lambda allows you to put code on the cloud, **without having to deal with running a server**. The code is **triggered** by a variety of **events**, and can perform basic tasks in response. Typically **stateless**.
- **Function** as a Service (FaaS)
- Examples of where Lambda might be useful:
 - IoT deployments Parsing incoming data from sensors
 - Automated backups An hourly script run and checking the backups!
 - ML deployment Refresh air pollution map every 10 minutes
 - Media conversion Managing the deployment of infrastructure, Netflix
 - Websites Low traffic sites with scripts

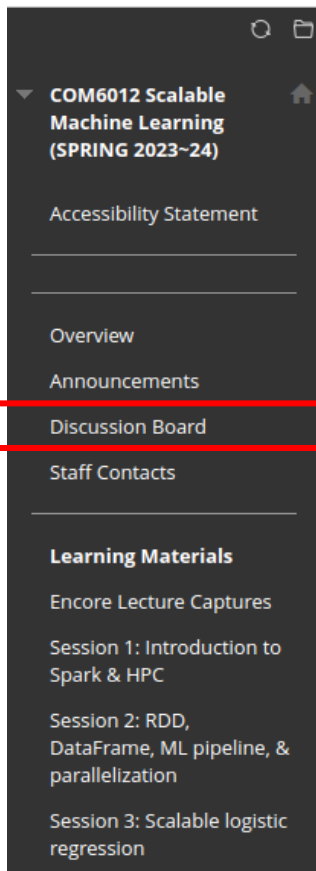
Activities

- AWS EC2 - Infrastructure as a Service (IaaS)
- AWS Lambda - Function as a Service (FaaS)

Take Home Messages

- Not always worth distributing your code (and lots of issues, including privacy).
- Service Models (Infrastructure-, Platform-/Function-, Software- as-a-Service).
 - EC2 is an example of Infrastructure-as-a-Service.
 - Lambda is an example of a Function-as-a-Service.
- Technology is changing quickly.
 - AWS EC2 is only about 15 years old.
 - Docker and Lambda are only about 10 years old.
 - Need to keep yourself updated if you enter this field.

Please note:



Assignment Brief:

<https://docs.google.com/document/d/1QSBkfnLLgf5qM0KWkeRayeAbrlTRKY4JZ8aYZA8npro/edit>

FAQ

<https://docs.google.com/document/d/1QSBkfnLLgf5qM0KWkeRayeAbrlTRKY4JZ8aYZA8npro/edit#bookmark=id.3ukytk47kosf>

You are required to complete all assignment questions using **batch mode**.

Acknowledgement

The slides for this lecture are adapted from [Dr Michael Smith](#)'s guest lecture slides in 2023.